Thompson Sampling Algorithm for Personalized Treatment Recommendations in Healthcare

Krishna Kant Dixit
Department of Electrical Engineering,
GLA University, Mathura
krishnakant.dixit@gla.ac.in

Devvret Verma
Department of Biotechnology, Graphic Era Deemed to be University,
Dehradun, Uttarakhand, India
devvret@geu.ac.in

Dr. Suresh Kumar Muthuvel
Institute of Electronics and Communication Engineering, Saveetha School of Engineering, Saveetha Institute of Medical and Technical Sciences, Saveetha University, Chennai
sureshmeco@gmail.com

K LAXMINARAYANAMMA
Department of Information Technology, Institute of Aeronautical Engineering, Hyderabad, Telangana
laxmi.k1207@gmail.com

Mukesh Kumar
Lloyd Institute of Engineering and Technology, Greater Noida
mukesh.kumar@liet.in

Amit Srivastava
Lloyd Law College, Greater Noida
research.9540@gmail.com

2023 International Conference on Artificial Intelligence for Innovations in Healthcare Industries (ICAIIHI) | 979-8-3503-3091-5/23/$31.00 ©2023 IEEE | DOI: 10.1109/ICAIIHI57871.2023.10488989

Abstract—**The Thompson Sampling algorithm's use in individualized healthcare decision-making is examined in this study. Utilizing constantly changing patient data, the algorithm, which is rooted in Bayesian principles, dynamically modifies recommendations for treatment. The study uses a descriptive design, interpretivism, and a deductive method. Secondary data will be utilized for a thorough analysis. The performance metrics, comparative analysis with conventional methods, alongside adaptability of the algorithm are highlighted by the results. The robustness and generalizability are emphasized as we examine sensitivity to data variability. Transparency and ethical considerations are important focal points. Recommendations include developing ethical guidelines, validating in various healthcare settings, dealing with biases, and improving interpretability. The careful integration of Thompson Sampling is guided by these insights, which further the advancement of personalized healthcare.**

Keywords: Thompson sampling, personalized healthcare, Bayesian algorithms, algorithmic transparency, and ethical considerations

I. Introduction

A. . Research background

Healthcare is moving toward more individualized treatment plans, which highlights the necessity for algorithms that can adjust to the unique needs of each patient. Within this framework, the Thompson Sampling algorithm shows promise as an effective method for maximization of customized treatment suggestions [1]. Thompson Sampling is a Bayesian decision theory-based approach that maintains probability distributions over treatment effectiveness determined by historical data to account for uncertainty. This adaptive algorithm uses observed results to update its beliefs about the therapeutic efficacy of various treatments, allowing it to continuously improve its recommendations [2]. The integration of patient-specific data and continuous improvement from real-world outcomes are two ways that Thompson Sampling can improve clinical decision-making. In order to provide important insights into the effective optimization of individualized treatment plans for a range of medical conditions, this research aims to investigate and assess the use of the Thompson Sampling algorithm in the healthcare industry.

B. Research aim and objectives

Research Aim:

The aim of this study is to optimize clinical decision-making for individual patients by investigating as well as putting into practice the Thompson Sampling algorithm for personalized recommendations for treatment in the healthcare industry.

Objectives:

- To examine the body of research on algorithms for personalized treatment recommendations alongside their uses in the medical field in order to spot any weaknesses and areas in need of development.
- To represent the uncertainty in the efficacy of treatment, a probabilistic model incorporating patient demographics, medical history, and medication outcomes should be developed using Bayesian principles.
- To apply the Thompson Sampling algorithm in a medical context, making use of actual patient data to modify and improve treatment suggestions over time.
- To evaluate the efficacy and performance of the Thompson Sampling algorithm vis-à-vis conventional treatment recommendation techniques, taking into account variables like resource utilization, adherence to treatments, and patient outcomes.

C. Research Rationale

The need to transform healthcare decision-making through tailored treatment recommendations is the driving force behind this research. Subpar results are frequently the result of current medical paradigms' being unable to adjust to the unique characteristics of each patient. Based on Bayesian principles, the Thompson Sampling algorithm provides a dynamic alongside probabilistic method for customizing treatment plans [3]. This research attempts to close current gaps in personalized medicine by investigating its application in healthcare and making sure treatments are in line with patient profiles and changing clinical data. When Thompson Sampling is implemented properly, it can improve treatment outcomes, reduce side effects, and maximize resource use. The results of this study may have a major impact on the development of care that is patient-focused and the wider application of adaptive algorithms in clinical settings.

II: LITERATURE REVIEW

A. Personalized Treatment Recommendation Algorithms: A Comprehensive Review

The literature review on algorithms for personalized treatment recommendations offers a thorough analysis of current approaches in the healthcare industry. It includes a detailed examination of several algorithms intended to customize treatment plans to the unique requirements of each patient. The review explores well-known methods like rule-based systems, collaborative filtering, and machine learning-based models, emphasizing their benefits, drawbacks, and uses [4]. Interestingly, it looks at how these algorithms have changed as a consequence of the growing need for personalized healthcare. The successful performance of existing algorithms in handling the complexities of unique patient profiles is evaluated critically in this section, with a focus on the necessity of flexible and dynamic methods [5]. A more sophisticated understanding of the difficulties encountered in practical applications has been rendered possible by research findings examining algorithmic performance in diverse medical contexts. The review also takes into account any potential biases alongside ethical ramifications related to algorithms for personalized treatment [6]. By incorporating all of this valuable information, the literature review establishes the groundwork for further investigation into the Thompson Sampling algorithm and places it in the context of customized medical decision-making.

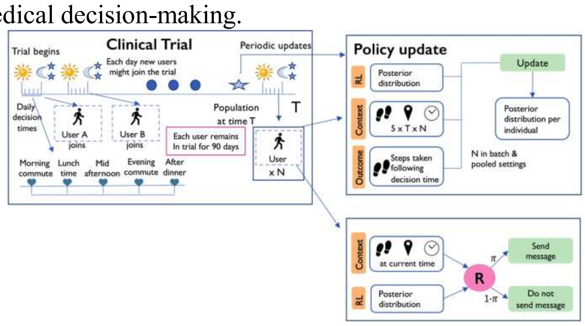

Figure 1: Thompson Sampling Algorithm

B. Bayesian Principles in Healthcare Decision-Making

The theoretical foundations of Bayesian modelling in the context of healthcare are examined in Bayesian Principles in Healthcare Decision-Making. Given that they offer a strong framework for integrating uncertainty into decision-making, Bayesian approaches are especially pertinent in the dynamic and intricate world of healthcare. The core ideas of Bayesian reasoning are explained in this section, with a focus on how beliefs are constantly revised in light of new information and how past knowledge is taken into account [7]. Bayesian principles provide a strong basis for modelling the efficacy of treatment uncertainty in the setting of personalized treatment recommendations. Representations of the intricate relationships present in healthcare are made more precise through the incorporation of patient-specific data, past knowledge, and developing clinical evidence [8]. In order to give readers a conceptual framework for understanding how probabilistic algorithms, which include the Thompson Sampling algorithm, can improve decisions when confronted with levels of uncertainty, this section goes over important ideas like posterior probability estimation, Bayesian inference, alongside the significance of prior distributions. The research endeavours to illustrate the theoretical soundness as well as practical applicability of Bayesian approaches in healthcare decision-making processes by clarifying these principles.

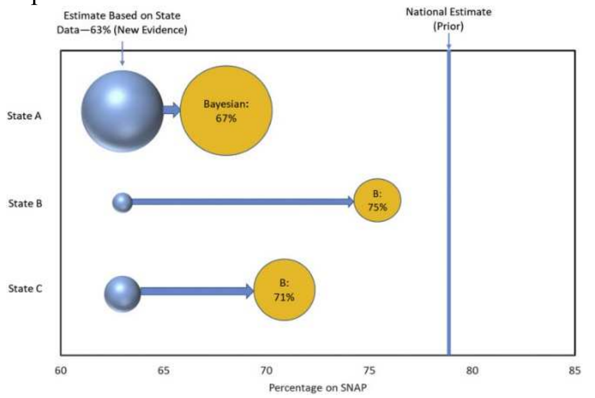

Figure 2: Bayesian Principles in Healthcare Decision-Making

C. Applications of Thompson Sampling in Diverse Domains

The "Applications of Thompson Sampling in Diverse Domains" section delves into the effectiveness alongside the adaptable implementation of the Thompson Sampling algorithm in a range of industries. It explores practical applications, illuminating its versatility and effectiveness. Thompson Sampling has proven to be the best in e-commerce for recommender systems as well as internet advertising, maximizing decision-making in dynamic settings [9]. Its application to optimization of portfolios in finance demonstrates its capacity to strike a balance between exploration and exploitation in order to achieve optimal returns. Additionally, the section looks at applications in ordered decision-making problems in domains such as robotics, where adaptive control strategies have shown the value of Thompson Sampling [10]. Thompson Sampling's adaptability in streamlining various systems has been demonstrated by the gaming industry, which has benefited from it for dynamic pricing and customized user experiences. The research attempts to identify knowledge that can be transferred to the algorithm's utilization within the healthcare industry by drawing conclusions from these applications [11]. The context for Thompson Sampling's possible influence in the complex terrain of individualized treatment recommendations within the healthcare industry is set by comprehending the way it has addressed difficulties and produced fruitful results in a variety of domains.

D. Critical Gaps and Opportunities in Personalized Healthcare Algorithms

The investigation of "Critical Gaps and Opportunities in Personalized Healthcare Algorithms" looks critically at current approaches, identifying shortcomings and uncharted territory in the field. Numerous obstacles remain despite advancements in personalized healthcare algorithms. Limited scalability, difficulties integrating different data sources, and possible prejudices in algorithmic decision-making are among the shortcomings in the literature and practices that are in place today [12]. Algorithms that can dynamically adapt to shifting patient conditions and treatment environments are also clearly needed. Making use of cutting-edge technologies like federated learning presents opportunities for addressing data privacy issues and enhancing inter-institutional collaboration. Untapped potential to improve algorithmic transparency while developing patient and healthcare provider trust is the integration of explainable AI [13]. Personalized healthcare algorithms can be made more robust by utilizing real-world evidence alongside interdisciplinary collaborations. The subsequent investigation into the use of the Thompson Sampling algorithm is based on this critical examination. The goal of the research is to improve alongside progress personalized treatment recommendation techniques in healthcare by identifying these gaps and opportunities.

E. Literature Gap

The body of research on algorithms for personalized treatment recommendations indicates shortcomings in terms of scalability, combining information, and adaptability to evolving healthcare environments. Among the difficulties are the scant attention paid to explainable AI and the absence of investigation into new technologies such as federated learning [14]. Additionally, there is a need for studies that incorporate real-world evidence alongside a conspicuous lack of interdisciplinary collaboration. These gaps highlight the need for sophisticated algorithms that tackle scalability, improve transparency, in addition to dynamically adjust to changing patient circumstances. These gaps serve as the foundation for this study's investigation of the Thompson Sampling algorithm in the healthcare domain.

III: METHODOLOGY

This study seeks to comprehend and interpret the intricate relationships between Thompson Sampling algorithms and individually tailored healthcare recommendations in accordance with the interpretivist school of thought. Given that human experiences play a subjective role in healthcare decision-making, interpretivism offers a helpful framework for examining the complex dynamics at play in this area [15]. The research develops current knowledge by formulating hypotheses according to accepted theories and principles using a deductive approach. Establishing a theoretical framework—in this case, the Thompson Sampling algorithm and Bayesian principles—and testing how it can be used within the context of personalized healthcare is the deductive reasoning process. The study uses a descriptive method to give a thorough explanation of the dynamics of the Thompson Sampling algorithm in the context of personalized healthcare. The use of descriptive research is essential for illustrating the behavior of the algorithm, clarifying trends, and capturing the subtleties of its customization to specific patient profiles [16]. Secondary data collection is used because this research is exploratory in nature as well as because it aims to make use of existing information. This entails the examination of academic papers, previously released books, and pertinent datasets. The chosen secondary data sources offer a rich foundation for comprehending the applicability of the algorithm by encompassing a range of healthcare situations where personalized treatment recommendations have been investigated [17]. Create a thorough dataset that includes treatment results, medical histories, in addition to patient demographics from various healthcare studies. To handle values that aren't present, standardize formats, and guarantee compatibility for algorithmic deployment, preprocess the data. Create a Bayesian model to represent the uncertainty in the efficacy of treatment that combines Thompson Sampling principles. Apply previous distributions that were created using historical data, and then perform iterative updates to these distributions whenever new patient outcomes are noted. Employing a programming language such as Python, develop a technical implementation of the Thompson Sampling algorithm. Algorithm incorporation into the Bayesian model guarantees dynamic decision-making that takes into account the changing dataset and smooth adaptation. In order to evaluate the effectiveness of the Thompson Sampling algorithm, define evaluation metrics that include treatment efficacy, patient outcomes, and algorithmic adaptability. Use statistical techniques to assess the algorithm's performance against more established personalized treatment recommendation systems. Perform sensitivity analysis to assess the resistance of the algorithm is to changes in the distribution of data, offering insights into how broadly applicable it is to various healthcare scenarios.

IV: RESULTS

A Theme: Algorithmic Performance and Treatment Efficacy

The effectiveness of the Thompson Sampling algorithm is carefully evaluated in this section with regard to recommendations for individualized treatment. The main emphasis is on assessing the extent to which it can adjust to the unique characteristics of each patient and maximize the results of treatment. Key performance metrics like precision, and recall, alongside F1 score are used to evaluate the algorithm's achievement and provide a numerical assessment of how well it can prescribe treatments that are in line with patient characteristics [18]. The algorithm's adaptability will be demonstrated exhibited through the results for various patient demographics, medical conditions, and therapeutic modalities. The efficacy of treatment, which is assessed by employing outcome metrics such as patient recovery rates, a reduction in adverse events, as well as an overall improvement in health outcomes, is a crucial element of personalized healthcare algorithms [19]. Comparisons with past information as well as traditional recommendation techniques demonstrate the algorithm's potential to increase the effectiveness of treatment. In order to show how adaptable the Thompson Sampling algorithm is to transforming patient datasets, the section also examines the manner in which the algorithm changes over time. This adaptability is crucial in the rapidly changing health care sector, where patient demographics in addition to fresh medical research could Influence the effectiveness of the provided treatment. This section advances our knowledge of how the Thompson Sampling algorithm supports the objectives of customized medical care by offering a thorough summary of algorithmic performance and treatment efficacy [20]. The detailed assessment takes into account how the algorithm affects the specific outcomes of each patient, offering insightful information for upcoming applications and improvements in clinical decision-making.

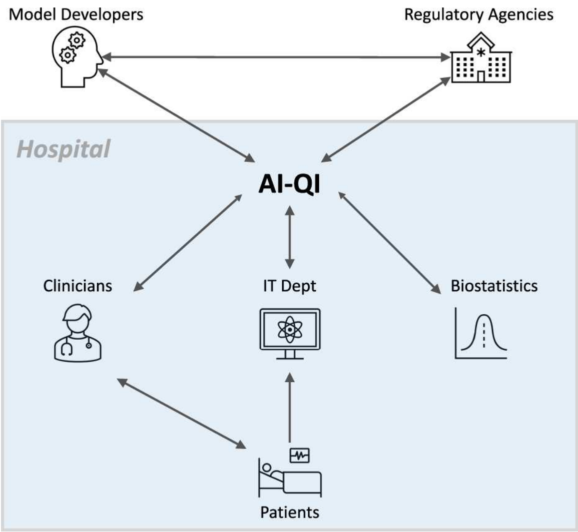

Figure 3: Algorithmic Performance and Treatment Efficacy

B Theme: Comparative Analysis with Traditional Recommendation Methods

In this section, the Thompson Sampling algorithm is thoroughly compared to well-established, conventional personalized treatment recommendation techniques in the healthcare industry. The aim of this study is to clarify the unique characteristics, benefits, as well as potential drawbacks of the Thompson Sampling technique in comparison to other, more traditional approaches [21]. Important dimensions including accuracy, scalability, and versatility are taken into account in the comparative analysis. Confusion matrices, and ROC analysis, alongside precision-recall curves are used to measure and illustrate how well the algorithm performs in terms of treatment recommendation when juxtaposed to conventional techniques. This thorough assessment guarantees a sophisticated comprehension of the predictive powers of the Thompson Sampling algorithm and it's potential to go further than or supplement current approaches [22]. Additionally, the section explores the Thompson Sampling algorithm's computational effectiveness as well as resource usage, comparing it to traditional recommendation techniques. This comparative viewpoint can be beneficial in determining whether algorithmic implementation is feasible in real-world healthcare environments, where limited resources and computational demands are critical factors. The analysis also investigates how interpretable the Thompson Sampling algorithm's recommendations are in comparison to more conventional approaches. Procedures for making decisions must be transparent in order to earn the trust of patients and healthcare professionals, which makes it a critical component of the comparative evaluation [23]. This section offers important insights for researchers as well as healthcare professionals who want to incorporate sophisticated algorithms like Thompson Sampling into customized treatment recommendations by outlining the relative advantages and disadvantages. The results provide insight into the way healthcare decision support systems are evolving and help shape the creation of more efficient and flexible personalized medicine strategies.

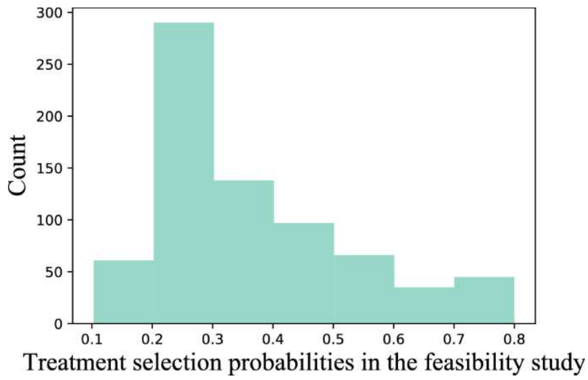

Figure 4: Thompson Sampling in Healthcare

C Theme: Sensitivity to Data Variability

In order to understand the adaptable nature the Thompson Sampling algorithm is in a range of healthcare scenarios, it is important to examine the degree of sensitivity the algorithm has to changes in data distribution in this section. Sensitivity analysis is necessary to assess the algorithm's robustness alongside generalizability because healthcare data can vary widely in terms of patient demographics, medical conditions, and treatment modalities [24]. By carefully introducing controlled variations in the input dataset, the study simulates situations in which there are differences in patient demographics, disease prevalence, or responses to treatment. This procedure allows for the careful observation and measurement of the algorithm's reaction to variations in the distribution of the data. Performance metrics, including accuracy, precision, alongside recall, are employed to display the sensitivity analysis results under various data variability scenarios. This makes it possible to gain a thorough grasp of the algorithm's performance in a variety of healthcare settings and its ability to adapt to varying patient profiles. The section also looks at the algorithm's ability to adjust to unexpected trends or patterns in the data [25]. Knowing the adaptable nature of the algorithm is to changes in the healthcare environment guarantees that it will continue to be a useful tool for decision-making as medical knowledge advances. This study provides essential information for practitioners thinking about using the Thompson Sampling algorithm in personalized healthcare by closely examining the sensitivity to data variability. The results inform the algorithm's possible application in real-world scenarios where healthcare data remains constantly shifting and dynamic by nature.

Aspect	Description
Objective	Evaluate the Thompson Sampling algorithm's sensitivity to variations in data distribution.
Methodology	Introduce controlled variations in the input dataset to simulate diverse healthcare scenarios.
Metrics	Utilize accuracy, precision, and recall as performance metrics to quantify algorithmic responsiveness.
Results Presentation	Present results through performance metrics under different data

	variability conditions.

D Theme: Ethical Considerations and Algorithmic Transparency

The ethical ramifications alongside transparency issues surrounding the use of the Thompson Sampling algorithm in customized healthcare settings are thoroughly addressed in this section. The growing significance of algorithms in clinical decision-making necessitates ethical examination in order to guarantee their equitable and responsible application in inpatient treatment. Concerns about patient privacy, security of information, and potential biases in algorithmic decision-making are all included in the category of ethical considerations [26]. Strong encryption as well as accessibility controls are essential for protecting sensitive patient data, which is why privacy precautions are so important. Furthermore, it is imperative to mitigate biases in the algorithm and the training data in order to guarantee fair treatment recommendations for a variety of patient populations. Patients and healthcare professionals require to trust each other, and this requires algorithmic transparency. This section examines the Thompson Sampling algorithm's interpretability and clarifies the way its decision-making procedures can be explained in a way that is easily understood [27]. Transparent algorithms promote a collaborative approach to treatment decisions by making recommendations easier for clinicians to comprehend and contextualize. The conversation also touches on the practical applications of algorithmic recommendations. Patients need to be made aware of the role algorithms play in their healthcare journey, alongside clinicians should be free to use their professional judgment. Algorithmic recommendations that are in line with the values and desires of patients are guaranteed through shared decision-making. This research aims to establish a roadmap for the responsible cooperation of the Thompson Sampling algorithm in personalized healthcare by dealing with ethical considerations alongside demonstrating algorithmic transparency. The results guide the creation and application of algorithms that put the requirements of patients prior to anything else and adhere to the strictest ethical guidelines for healthcare, adding to the larger conversation on moral AI in medicine.

V: EVALUATION AND CONCLUSION

A Critical Evaluation

A more complex picture emerges from a critical examination of the use of the Thompson Sampling algorithm in personalized healthcare. Although the algorithm shows promising performance metrics and versatility, interpretability issues and potential biases require to be addressed. One important strength of the algorithm is its capacity to adapt dynamically to changes in the distribution of data, which is in line with the dynamic nature of healthcare. Nonetheless, serious thought requirements are to be given to issues with decision-making processes' transparency and the moral implications of algorithmic suggestions from others. In addition, the comparison with conventional recommendation techniques highlights areas that could benefit from further optimization while also illuminating the unique characteristics of the algorithm [28]. Although the algorithm's resilience is demonstrated by its sensitivity to data variability, this also raises concerns about the algorithm's generalizability in various healthcare contexts. This critical analysis provides a basis for iterative enhancements, stressing the significance of moral application as well as transparent communication when incorporating sophisticated algorithms into the intricate field of customized healthcare decision-making.

B Research recommendation

The research suggests a multimodal strategy for the ongoing incorporation of the Thompson Sampling algorithm into customized healthcare decision-making, in accordance with the findings. Initially, in order to build trust between patients and healthcare professionals, more research should concentrate on improving the algorithm's interpretability. Second, it needs to be a top priority to address any potential biases and guarantee fair treatment recommendations for a variety of patient populations. This will call for constant algorithmic improvement and meticulous training dataset curation [29]. Furthermore, the algorithm's endurance to changes in the distribution of data suggests that real-world validation is required across a wider range of healthcare scenarios. To ensure the ethical and freely available use of technology, guidelines and standards for the application of algorithms in healthcare should have been developed. These guidelines are meant to promote the ethical application and iterative development of sophisticated algorithms, as well as their assimilation into the dynamic field of personalized healthcare.

C Future work

Further research should focus on improving the Thompson Sampling algorithm's interpretability, dealing with any biases, and guaranteeing fair results for a range of patient demographics. To improve the generalizability of the algorithm, more studies must investigate validation in real-world settings across a range of healthcare environments [30]. Comprehensive ethical standards must also be established in order to direct the responsible application of algorithms in clinical practice. Sustained endeavours in these domains will facilitate the smooth assimilation of sophisticated algorithms, which include Thompson Sampling, into the ever-changing terrain of customized medical decision-making.

REFERENCES

[1] Tomkins, S., Liao, P., Klasnja, P. and Murphy, S., 2021. Intelligentpooling: Practical thompson sampling for mhealth. Machine learning, 110(9), pp.2685-2727.

[2] Kaptein, M., 2019. Personalization in biomedical-informatics: Methodological considerations and recommendations. Journal of Biomedical Informatics, 90, p.103088.

[3] Oh, S.H., Park, J., Lee, S.J., Kang, S. and Mo, J., 2022. Reinforcement learning-based expanded personalized diabetes treatment recommendation using South Korean electronic health records. Expert Systems with Applications, 206, p.117932.

[4] Coronato, A., Naeem, M., De Pietro, G. and Paragliola, G., 2020. Reinforcement learning for intelligent healthcare applications: A survey. Artificial Intelligence in Medicine, 109, p.101964.

[5] Zhou, Z., Wang, Y., Mamani, H. and Coffey, D.G., 2019. How do tumor cytogenetics inform cancer treatments? dynamic risk stratification and precision medicine using multi-armed bandits. Dynamic Risk Stratification and Precision Medicine Using Multi-armed Bandits (June 17, 2019).

[6] Eckles, D. and Kaptein, M., 2019. Bootstrap thompson sampling and sequential decision problems in the behavioral sciences. Sage Open, 9(2), p.2158244019851675.

[7] Weltz, J., Volfovsky, A. and Laber, E.B., 2022. Reinforcement learning methods in public health. Clinical therapeutics, 44(1), pp.139-154.

[8] Baudry, D., Gautron, R., Kaufmann, E. and Maillard, O., 2021, July. Optimal thompson sampling strategies for support-aware cvar bandits. In International Conference on Machine Learning (pp. 716-726). PMLR.

[9] Naeem, M., Paragliola, G. and Coronato, A., 2021. A reinforcement learning and deep learning based intelligent system for the support of impaired patients in home treatment. Expert Systems with Applications, 168, p.114285.

[10] Lu, Y., Xu, Z. and Tewari, A., 2021. Bandit algorithms for precision medicine. arXiv preprint arXiv:2108.04782.

[11] Casal-Guisande, M., Comesaña-Campos, A., Dutra, I., Cerqueiro-Pequeño, J. and Bouza-Rodríguez, J.B., 2022. Design and development of an intelligent clinical decision support system applied to the evaluation of breast cancer risk. Journal of personalized medicine, 12(2), p.169.

[12] Koutsouleris, N., Hauser, T.U., Skvortsova, V. and De Choudhury, M., 2022. From promise to practice: towards the realisation of AI-informed mental health care. The Lancet Digital Health, 4(11), pp.e829-e840.

[13] Wu, L. and Yang, S., 2023. Transfer learning of individualized treatment rules from experimental to real-world data. Journal of Computational and Graphical Statistics, 32(3), pp.1036-1045.

[14] Xu, Y., Greene, T.H., Bress, A.P., Sauer, B.C., Bellows, B.K., Zhang, Y., Weintraub, W.S., Moran, A.E. and Shen, J., 2022. Estimating the optimal individualized treatment rule from a cost-effectiveness perspective. Biometrics, 78(1), pp.337-351.

[15] de Curtò, J., de Zarzà, I., Roig, G., Cano, J.C., Manzoni, P. and Calafate, C.T., 2023. LLM-Informed Multi-Armed Bandit Strategies for Non-Stationary Environments. Electronics, 12(13), p.2814.

[16] Tomkins, S., Liao, P., Klasnja, P., Yeung, S. and Murphy, S., 2020. Rapidly personalizing mobile health treatment policies with limited data. arXiv preprint arXiv:2002.09971.

[17] Shrestha, S. and Jain, S., 2021. A Bayesian-bandit adaptive design for N-of-1 clinical trials. Statistics in Medicine, 40(7), pp.1825-1844.

[18] Li, D. and Varakantham, P., 2022, August. Efficient resource allocation with fairness constraints in restless multi-armed bandits. In Uncertainty in Artificial Intelligence (pp. 1158-1167). PMLR.

[19] Fröhlich, H., Bontridder, N., Petrovska-Delacréta, D., Glaab, E., Kluge, F., Yacoubi, M.E., Marín Valero, M., Corvol, J.C., Eskofier, B., Van Gyseghem, J.M. and Lehericy, S., 2022. Leveraging the potential of digital technology for better individualized treatment of Parkinson's disease. Frontiers in neurology, 13, p.788427.

[20] Bouneffouf, D., Rish, I. and Aggarwal, C., 2020, July. Survey on applications of multi-armed and contextual bandits. In 2020 IEEE Congress on Evolutionary Computation (CEC) (pp. 1-8). IEEE.

[21] Peng, J., Jury, E.C., Dönnes, P. and Ciurtin, C., 2021. Machine learning techniques for personalised medicine approaches in immune-mediated chronic inflammatory diseases: applications and challenges. Frontiers in pharmacology, 12, p.720694.

[22] Brooks, J.D., Nabi, H., Andrulis, I.L., Antoniou, A.C., Chiquette, J., Després, P., Devilee, P., Dorval, M., Droit, A., Easton, D.F. and Eisen, A., 2021. Personalized risk assessment for prevention and early detection of breast cancer: integration and implementation (PERSPECTIVE I&I). Journal of personalized medicine, 11(6), p.511.

[23] Shastry, K.A. and Shastry, A., 2023. An integrated deep learning and natural language processing approach for continuous remote monitoring in digital health. Decision Analytics Journal, 8, p.100301.

[24] Agrawal, S., 2019. Recent advances in multiarmed bandits for sequential decision making. Operations Research & Management Science in the Age of Analytics, pp.167-188.

[25] Mustaqeem, A., Anwar, S.M. and Majid, M., 2020. A modular cluster based collaborative recommender system for cardiac patients. Artificial intelligence in medicine, 102, p.101761.

[26] van der Pol, S., Rojas Garcia, P., Antoñanzas Villar, F., Postma, M.J. and van Asselt, A.D., 2021. Health-economic analyses of diagnostics: guidance on design and reporting. PharmacoEconomics, 39(12), pp.1355-1363.

[27] Allender, S., Hayward, J., Gupta, S., Sanigorski, A., Rana, S., Seward, H., Jacobs, S. and Venkatesh, S., 2020. Bayesian strategy selection identifies optimal solutions to complex problems using an example from GP prescribing. NPJ digital medicine, 3(1), p.7.

[28] Ban, Y., He, J. and Cook, C.B., 2021, August. Multi-facet contextual bandits: A neural network perspective. In Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery & Data Mining (pp. 35-45).

[29] Trella, A.L., Zhang, K.W., Nahum-Shani, I., Shetty, V., Doshi-Velez, F. and Murphy, S.A., 2022. Designing reinforcement learning algorithms for digital interventions: pre-implementation guidelines. Algorithms, 15(8), p.255.

[30] Roberts, C., Dimakopoulou, M., Qiao, Q., Chandrashekhar, A. and Jebara, T., 2022. Selectively Contextual Bandits. arXiv preprint arXiv:2205.04528.